




Matrix Data Visualization

Milan Vojnovic

ST445 Managing and Visualizing Data

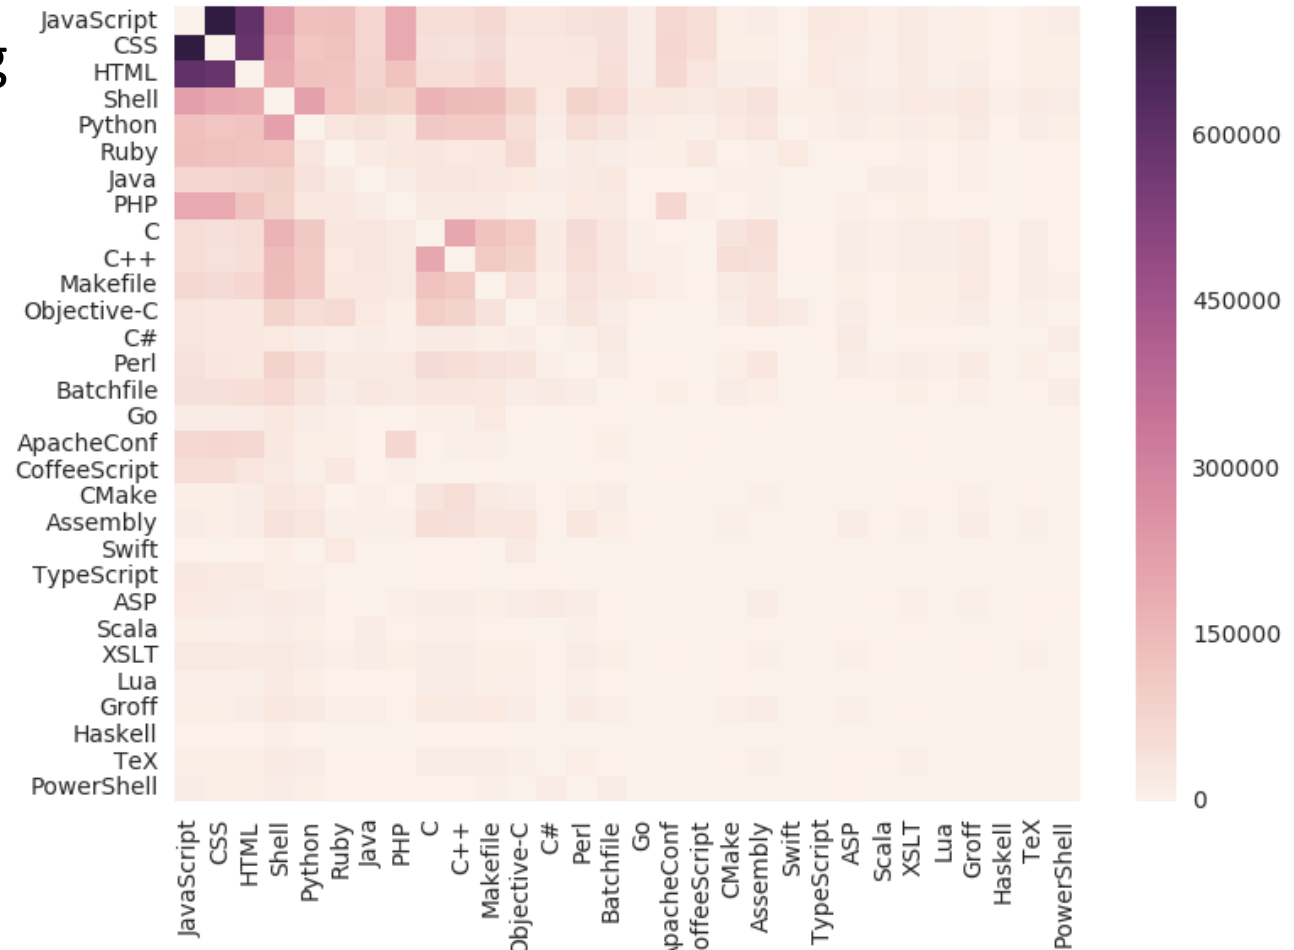
GitHub archive dataset

- The plot on this slide shows the percentage of repositories that use specific programming languages
- Suppose our goal is to visualize co-occurrence of pairs of programming languages in different repositories



Heatmap

- Co-occurrence of pairs of programming languages in repositories can be visualized by a heatmap
- It is important how the rows and columns are sorted to visualize any possibly existing clusters
- In this slide, the rows and columns are sorted in decreasing popularity of programming languages
 - A heuristic that for this instance reveals some clusters, but this is not necessary in general

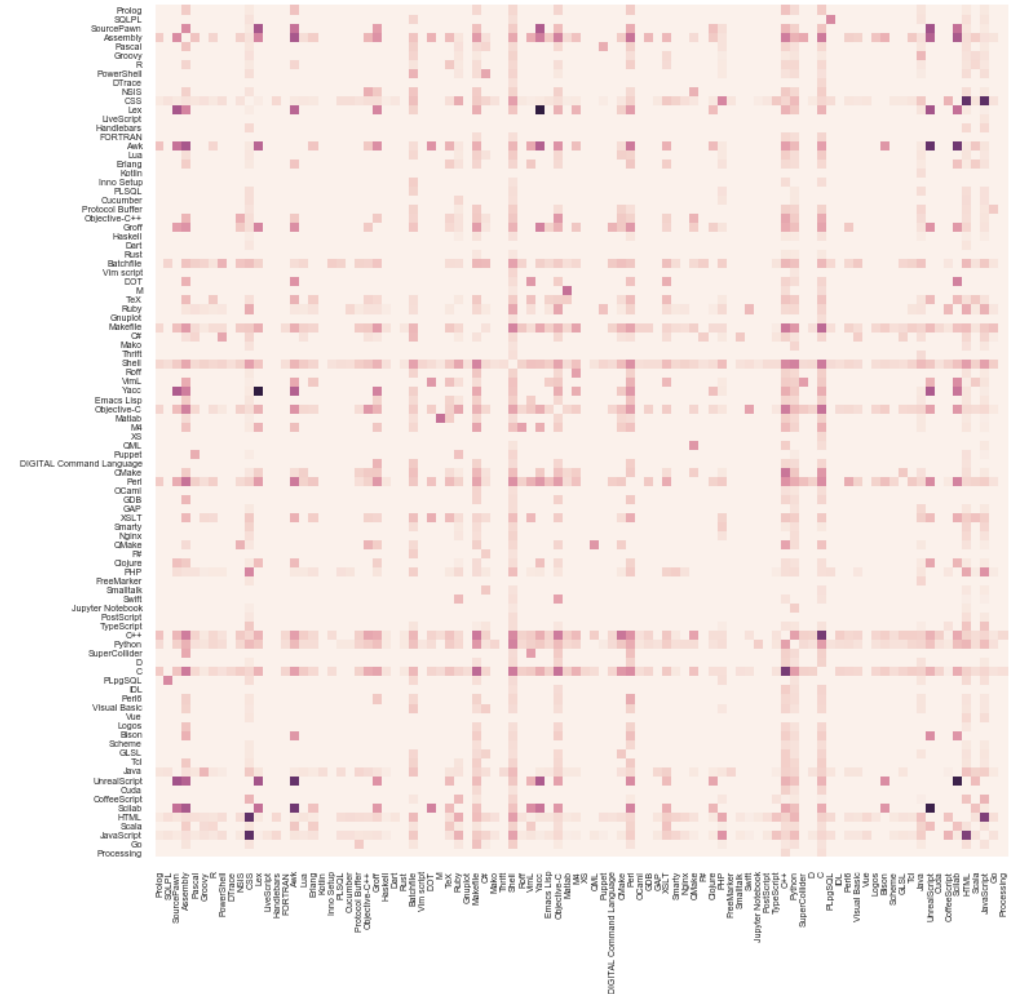


Use case: similarity matrices

- A similarity matrix quantifies how “similar” pairs of items are
- Similarity may be defined as **cosine similarity** for items associated with feature vectors: for two non-null feature vectors a and b in \mathbf{R}^n , the cosine similarity is

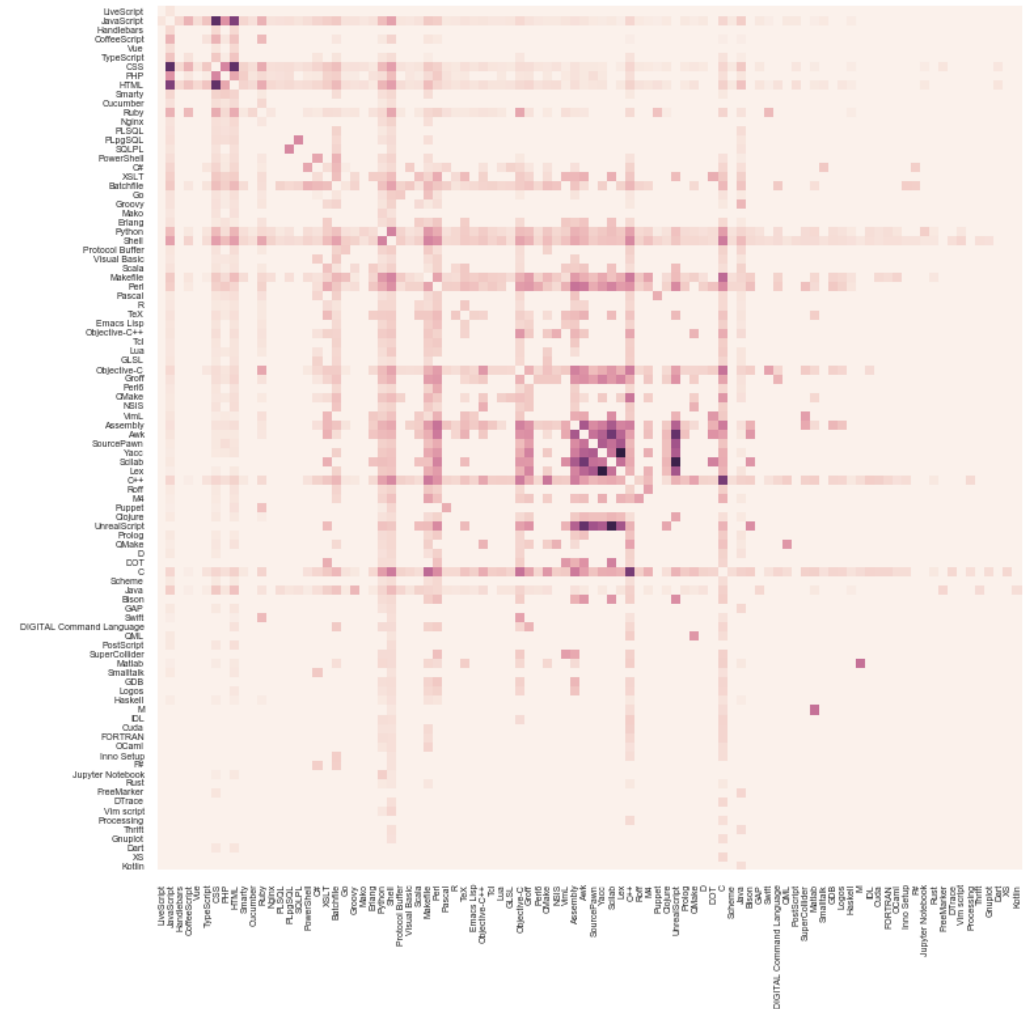
$$\text{sim}(a, b) = \frac{\sum_{i=1}^n a_i b_i}{\sqrt{\sum_{i=1}^n a_i^2} \sqrt{\sum_{i=1}^n b_i^2}}$$

- In our example, a feature vector associated with a programming language indicates its usage over different repositories



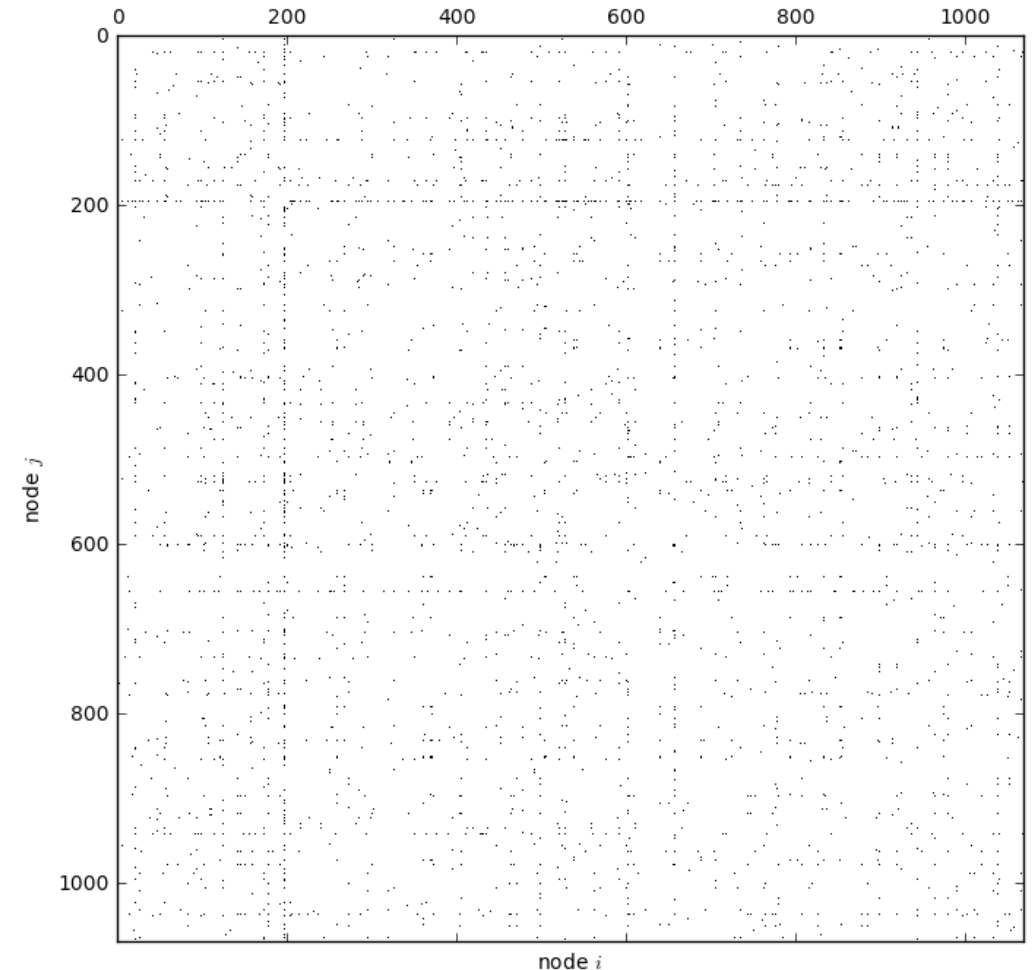
Use case: similarity matrices (cont'd)

- How can we order rows and columns of a matrix to visualize any possibly existing clusters in matrix data?



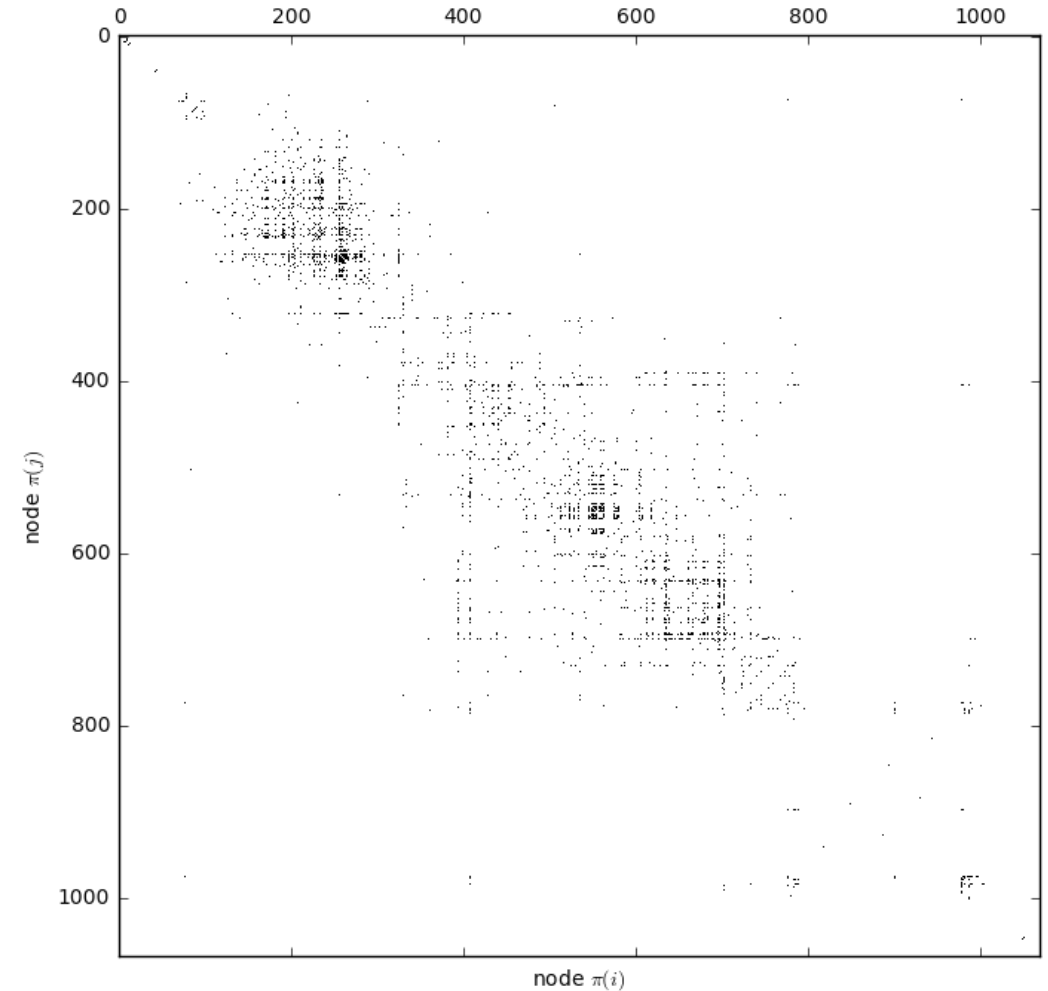
Use case: adjacency matrices

- Adjacency matrix of a graph indicates whether or not there is an edge between different pairs of nodes of the graph
- The plot in this slide shows the adjacency matrix of a graph that specifies existence of communications between Amazon Mechanical Turk workers
- Dataset source: Yin et al, The Communication Network within the Crowd, WWW 2016



Use case: adjacency matrices

- Reordered rows and columns
i.e. rearranged node identifiers
- The adjacency matrix with reordered rows and columns in this slide suggests existence of different communities of workers



Software modules

- Scikit-learn biclustering module
 - Scikit-learn Section 2.4 Biclustering

<http://scikit-learn.org/stable/modules/biclustering.html>

- from sklearn.cluster.bicluster import SpectralCoclustering
- from sklearn.cluster.bicluster import SpectralBiclustering

- R seriation package

- <https://cran.r-project.org/web/packages/seriation/index.html>

Some linear algebra concepts

Eigenvalues and eigenvectors

- λ is an **eigenvalue** of matrix A if for some vector $x \neq 0$

$$Ax = \lambda x$$

- A corresponding vector x is called an **eigenvector**

Laplacian matrix

- Let A be a real symmetric matrix
- The Laplacian matrix L_A is defined by

$$L_A = D_A - A$$

where D_A is a diagonal matrix with $d_{i,i} = \sum_{j=1}^n a_{i,j}$

- A is a real symmetric matrix $\Rightarrow L_A$ is a real symmetric matrix

Laplacian matrix (cont'd)

- For any real symmetric matrix $A \in \mathbf{R}^{n \times n}$:
 - A has n eigenvectors
 - All eigenvectors of A are pairwise orthogonal
 - All eigenvalues of A are real
- Every Laplacian matrix has all its eigenvalues real and non-negative, which is equivalent to saying that it is **positive semi-definite**
- Every Laplacian matrix has the vector of all ones e as an eigenvector corresponding to the eigenvalue zero

Fiedler value and eigenvector

- For real symmetric matrix $A \in \mathbf{R}^{n \times n}$, **Fiedler value** is defined as the minimum eigenvalue of the Laplacian L_A that has an eigenvector orthogonal to e
- The corresponding eigenvector is called a **Fiedler eigenvector**
- The Fiedler value is the optimum value of the optimization problem:

$$\begin{array}{ll} \text{minimize} & x^T L_A x \\ \text{subject to} & x^T e = 0 \\ & x^T x = 1 \end{array}$$

Historical remarks

- Miroslav Fiedler
- Czech Republic's mathematician
- 1926-2015
- Charles University, Prague



- M. Fiedler, Algebraic connectivity of graphs, Czechoslovak Math. J., 23(98), 1973

Laplacian matrix (cont'd)

- **Lemma:** For any real, symmetric matrix A :

$$x^T L_A x = \sum_{i < j} a_{i,j} (x_i - x_j)^2$$

- Laplacian matrices are closely related to **graph cuts**
 - Graph $G = (V, E)$ with edge weights: edge (i, j) has weight $a_{i,j}$
 - A is the adjacency matrix of G
 - Let x takes value in $\{-1, 1\}^n$ defining a vertex cut: partitioning vertices into two sets (negative and positive labeled)
 - **Graph cut** is defined as the sum of weights of edges whose end vertices belong to different components of the vertex cut

Seriation

Seriation

- Input: a real, symmetric matrix $A \in \mathbf{R}^{n \times n}$
 - $a_{i,j}$ interpreted as the similarity between items i and j
- Goal: find a linear ordering (permutation) of items such that similar items are placed nearby, specifically, find a permutation π^* that minimizes

$$c(\pi) = \sum_{i < j} a_{i,j} (\pi_i - \pi_j)^2$$

over the set of all possible permutations of n elements Π_n

- This problem is NP hard

Fractional relaxation

- Find optimal solution to the following problem:

$$\begin{array}{ll}\text{minimize} & c(x) \\ \text{subject to} & x^T e = 0 \\ & x^T x = 1 \\ & x \in \mathbf{R}^n\end{array}$$

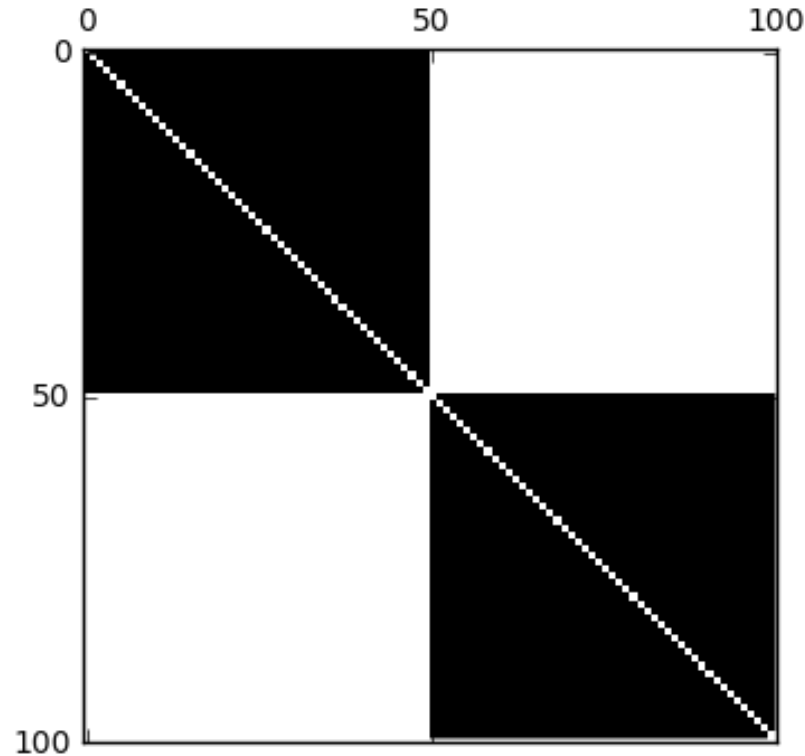
- The first constraint avoids multiplicity of solutions by adding a positive constant to each coordinate of x
- The second constraint avoids vector 0 to be a trivial solution
- The optimum solution is a Fiedler vector of A

Robinson matrices

- A matrix A is said to be a **Robinson matrix (R-matrix)** if, and only if,
 - A is symmetric
 - $a_{i,j} \leq a_{i,k}$ for $j < k < i$ and $a_{i,j} \geq a_{i,k}$ for $i < j < k$
- In other words, a matrix is a R-matrix if it is symmetric and it has off-diagonal elements non-decreasing by moving away from the diagonal
- A matrix A is said to be a **pre-R matrix** if it can be symmetrically permuted (reordering rows and columns by the same permutation) to become an R-matrix

An example of a R-matrix

$$A = (a_{i,j})$$



- Clearly, a symmetric matrix
- The elements are decreasing as we move away from the diagonal along any row

black point if $a_{i,j} = 1$
white point if $a_{i,j} = 0$

An example R-matrix: stochastic block model

- Stochastic block model is a random graph model commonly studied in the community detection and graph clustering literature
- A stochastic block model can be defined by
 - Parameters $0 \leq q \leq p \leq 1$
 - Set of vertices $V = \{1, 2, \dots, n\}$
 - Hidden bipartition of vertices $(S, V \setminus S)$
 - Set of edges E : for every pair of vertices (i, j) such that $i < j$ we have

$$(i, j) \in E \text{ independently with probability} = \begin{cases} p & \text{if } i, j \in S \text{ or } i, j \in V \setminus S \\ q & \text{otherwise} \end{cases}$$

- A is defined as the adjacency matrix of G
- The case $q = 0$ and $p = 1$ is trivial when the vertex-set partition is not hidden

Fiedler vector of R-matrices

- **Theorem:** If A is a R-matrix then it has a monotone Fiedler vector.
- Proof: Atkins et al (1998)

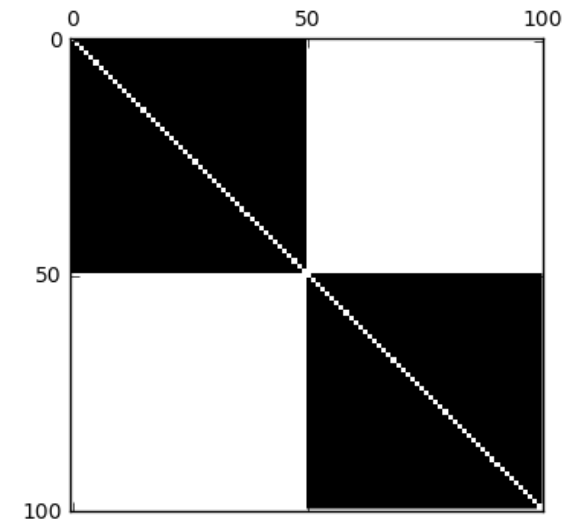
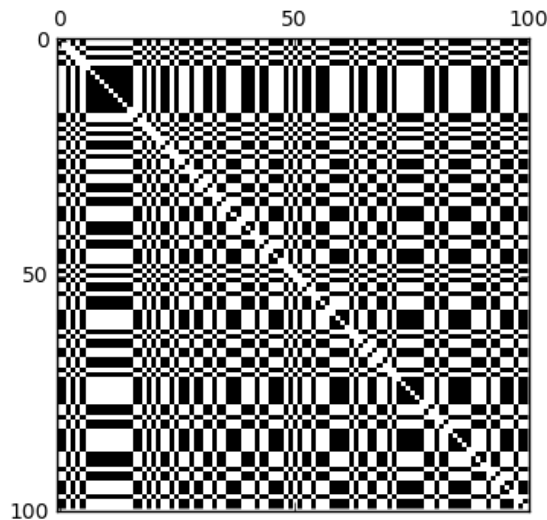
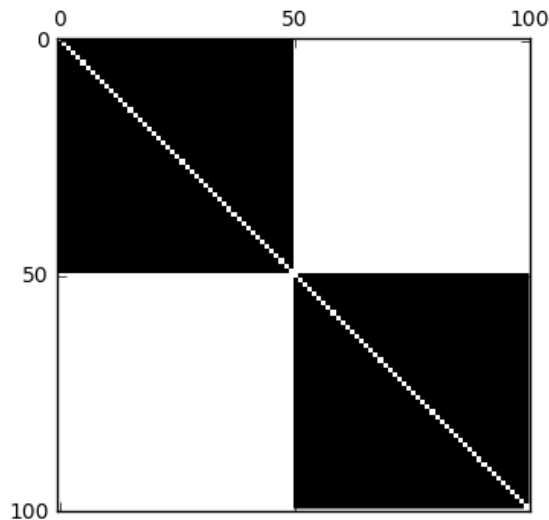
Fiedler vector of R-matrices (cont'd)

- **Theorem:** Let A be a pre-R matrix with a **simple** Fiedler value (i.e. unique value) and a Fiedler vector with no repeated elements. Let π_1 be the permutation sorting elements of the Fiedler vector in increasing order. Let π_2 be the permutation sorting elements of the Fiedler vector in decreasing order. Let Π_1 and Π_2 be the corresponding permutation matrices.

Then, $\Pi_1 A \Pi_1$ and $\Pi_2 A \Pi_2$ are R-matrices and **not other symmetric permutation of A produces an R matrix.**

- Proof: Atkins et al (1998)
- See Thm 4.7 in Atkins et al for a statement under weaker assumptions

Example 1: $p = 1, q = 0$

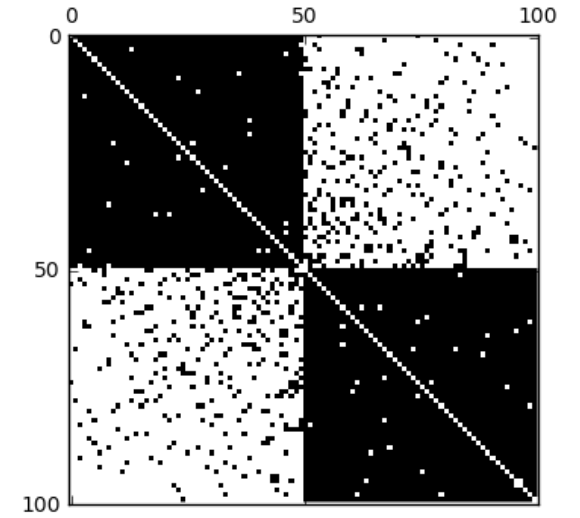
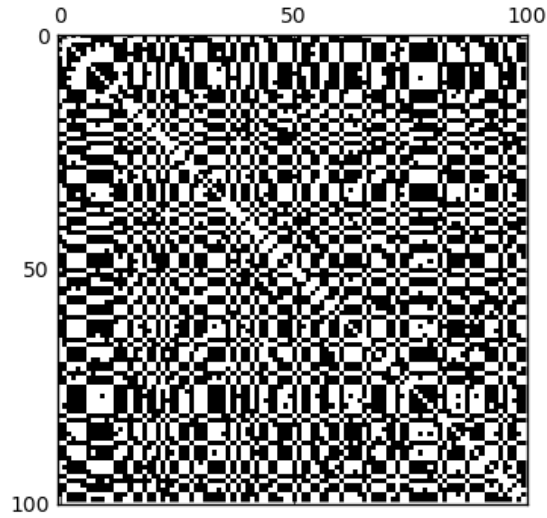
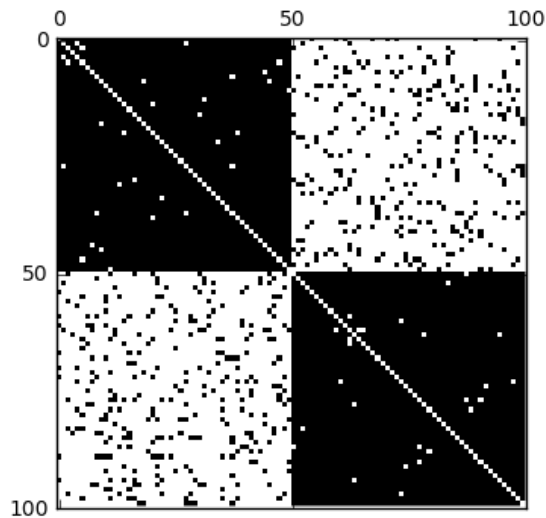


Rows and columns
permuted by the same
random permutation

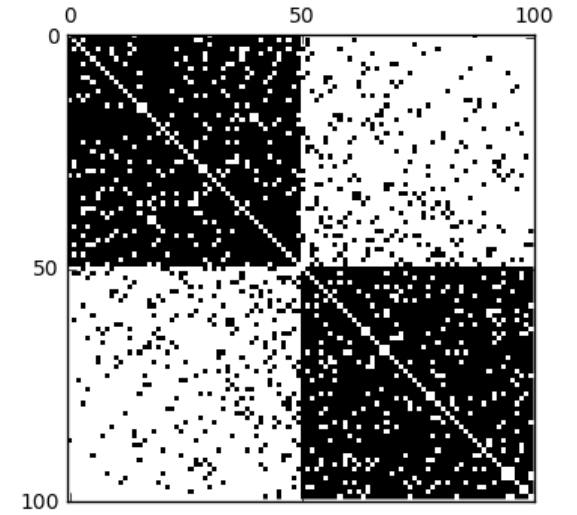
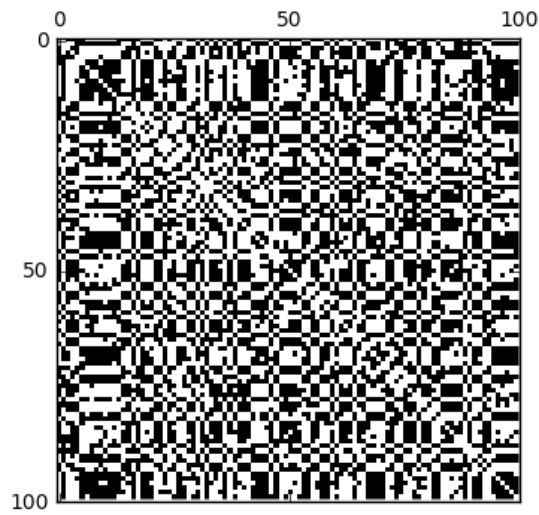
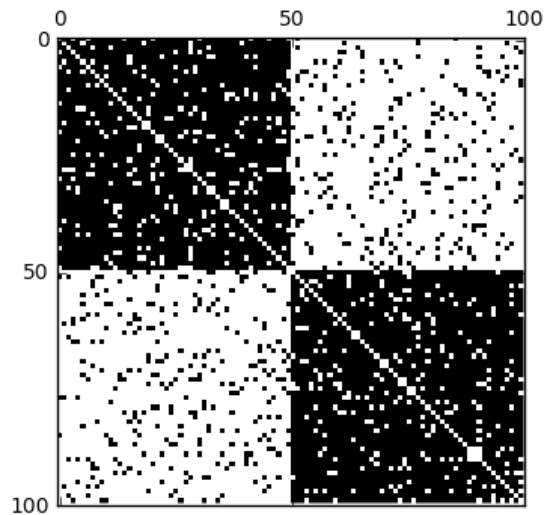
“symmetric permutation”

Rows and columns
sorted in a monotonic
order of the Fiedler
vector elements

Example 2: $p = 0.99$, $q = 0.01$



Example 3: $p = 0.9$, $q = 0.1$



Spectral co-clustering

Graph cuts

- Let V be a set of vertices with weights $w_{i,j}$ for $i, j \in V$
- **k-way partition**: let $P_k(S)$ be the set of all possible partitions of S in k components
 - In particular, $(S_1, S_2) \in P_2(S)$ is referred as a **bipartition**
- The **cut function** is defined as the sum of weights of cut edges:

$$\text{cut}(V_1, V_2) = \sum_{i \in V_1, j \in V_2} w_{i,j}$$

- More generally, we define

$$\text{cut}(V_1, V_2, \dots, V_k) = \sum_{1 \leq i < j \leq k} \text{cut}(V_i, V_j)$$

Graph cuts and the Laplacian matrix

- A bipartition (V_1, V_2) can be represented by a **partition vector** x defined as

$$x_i = \begin{cases} 1 & \text{if } i \in V_1 \\ -1 & \text{if } i \in V_2 \end{cases}$$

- **Lemma:** Any bipartition (V_1, V_2) and the corresponding partition vector x satisfy

$$\frac{x^T L_A x}{x^T x} = \frac{1}{n} 4 \text{ cut}(V_1, V_2)$$

Normalized cut function

- Define the weight of a set of vertices S : $w(S) = \sum_{i \in S} w_i$
- The **normalized cut function** is defined as follows:

$$Q(V_1, V_2) = \frac{\text{cut}(V_1, V_2)}{w(V_1)} + \frac{\text{cut}(V_1, V_2)}{w(V_2)}$$

- This cut function captures *both* **sparsity of edge cuts** and **balancing of the component sizes**
- Optimization problem formulation: minimize $Q(V_1, V_2)$ over $P_2(V)$

Special normalized cut functions

- **Ratio-cut**

- Each vertex has a unit weight
- Amounts to looking at the eigenvalue problem: $L_A y = \lambda y$

- **Normalized-cut**

- Each vertex weight is equal to the sum of weights of incident edges
- In this case $w(V_i) = \text{cut}(V_1, V_2) + w^{int}(V_i)$
where $w^{int}(V_i) = \sum_{(u,v) \in E: u,v \in V_i} w_{u,v}$
- Amounts to the generalized eigenvalue problem : $L_A y = \lambda D_A y$

Normalized-cut function

- Show that

$$Q(V_1, V_2) = 2 - \left(\frac{w^{int}(V_1)}{w(V_1)} + \frac{w^{int}(V_2)}{w(V_2)} \right)$$

- Minimizing the normalized-cut is equivalent to maximizing the proportion of edge weights that lie within each component of a vertex set partition

Generalized partition-vector representation

- Let y be a **generalized partition vector** defined as

$$y_i = \begin{cases} \sqrt{\frac{w(V_2)}{w(V_1)}} & \text{if } i \in V_1 \\ -\sqrt{\frac{w(V_1)}{w(V_2)}} & \text{if } i \in V_2 \end{cases}$$

- **Properties:** (a) $y^T D_w y = w(V)$ and (b) $y^T D_w e = 0$

- **Lemma:**

$$\frac{y^T L_A y}{y^T D_w y} = \frac{\text{cut}(V_1, V_2)}{w(V_1)} + \frac{\text{cut}(V_1, V_2)}{w(V_2)}$$

Fractional relaxation

$$\text{minimize } \frac{y^T L_A y}{y^T D_w y}$$

$$\text{subject to } y^T D_w e = 0$$

$$y \neq 0$$

$$y \in \mathbf{R}^n$$

- **Theorem:** The solution is the eigenvector corresponding to the second smallest eigenvalue λ_2 of the generalized eigenvalue problem:

$$L_A y = \lambda D_w y$$

- **Corollary:** The optimum value of the min normalized cut is $\geq \lambda_2$

Bipartite graph clustering

- Bipartite graph
 - The vertex set consists of disjoint sets of vertices L and R
 - Each edge has its vertices in L and R
- Let D_L and D_R be diagonal matrices with diagonal elements

$$(D_L)_{i,i} = \sum_{j \in R} w_{i,j} \text{ and } (D_R)_{i,i} = \sum_{j \in L} w_{j,i}$$

- For graph G :

$$A = \begin{pmatrix} 0 & W \\ W^T & 0 \end{pmatrix}, D_A = \begin{pmatrix} D_L & 0 \\ 0 & D_R \end{pmatrix} \text{ and } L_A = \begin{pmatrix} D_L & -W \\ -W^T & D_R \end{pmatrix}$$

The generalized eigenvalue problem

- The generalized eigenvalue problem can be written as

$$\begin{aligned} D_L x - W y &= \lambda D_L x \\ -A^T x + D_R y &= \lambda D_R y \end{aligned}$$

- **Assumption:** W has a strictly positive element in each row and column
- Note that we can write:

$$\begin{aligned} D_L^{1/2} x - D_L^{-1/2} W y &= \lambda D_L^{1/2} x \\ -D_R^{-1/2} A^T x + D_R^{1/2} y &= \lambda D_R^{1/2} y \end{aligned}$$

Change of variables

- Using the change of variables

$$\tilde{W} = D_L^{-1/2} W D_R^{-1/2}$$

$$u = D_L^{1/2} x$$

$$v = D_R^{1/2} y \text{ and}$$

$$\sigma = 1 - \lambda$$

we can write

$$\tilde{W} v = \sigma u \quad \text{and} \quad \tilde{W}^T u = \sigma v$$

Eigenvectors

- The eigenvector x_2 corresponding to the second smallest eigenvalue λ_2 of the generalized eigenvalue problem can be written as:

$$x_2 = \begin{pmatrix} D_L^{-1/2} u_2 \\ D_R^{-1/2} v_2 \end{pmatrix}$$

where u_2 and v_2 are the left and right singular vectors of \tilde{W} corresponding to the singular value $\sigma_2 = 1 - \lambda_2$

- We can think of u_2 to give a partition of the set of left vertices and v_2 to give a partition of the set of right vertices

Bi-clustering algorithm

- Input: W
- Compute $\tilde{W} = D_L^{-1/2} W D_R^{-1/2}$
- Compute the left and right singular vectors u_2 and v_2 of \tilde{W} corresponding to the singular value $\sigma_2 = 1 - \lambda_2$
- Partition the set of vertices in two components using k -means algorithm for input data points x_2

k-way clustering algorithm

- Given a positive integer $k \geq 2$ the goal is to partition the set of left vertices and the set of right vertices in k components
- Let $U = (u_2, u_3, \dots, u_{\ell+1})$ and $V = (v_2, v_3, \dots, v_{\ell+1})$ be ℓ the left and right singular vectors
- Let $(x_2, x_3, \dots, x_{\ell+1}) = \begin{pmatrix} D_L^{-1/2} U \\ D_R^{-1/2} V \end{pmatrix}$
- Apply the k-means algorithm to the input ℓ -dimensional points

Evaluating a bi-clustering: consensus score

- The definition of the consensus score as implemented in `sklearn.metrics.consensus_score`
- Quantifies similarity between two input *sets* of biclusters
- Similarity between individual biclusters is computed
 - Default is Jaccard similarity for two sets A and B defined as $|A \cap B| / |A \cup B|$
- The best matching between sets is found using the Hungarian algorithm
- The final score is the sum of the similarities divided by the size of the larger set

```
def consensus_score(a, b, similarity="jaccard"):
```

Parameters

a : (rows, columns)

 Tuple of row and column indicators for a set of biclusters.

b : (rows, columns)

 Another set of biclusters like ``a``.

similarity : string or function, optional, default: "jaccard"

 May be the string "jaccard" to use the Jaccard coefficient, or
 any function that takes four arguments, each of which is a 1d
 indicator vector: (a_rows, a_columns, b_rows, b_columns).

```
if similarity == "jaccard":
```

```
    similarity = _jaccard
```

```
matrix = _pairwise_similarity(a, b, similarity)
```

```
indices = linear_assignment(1. - matrix)    # maximum weight matching (input is a cost matrix)
```

```
n_a = len(a[0])                            from sklearn.utils.linear_assignment_
```

```
n_b = len(b[0])
```

```
return matrix[indices[:, 0], indices[:, 1]].sum() / max(n_a, n_b)
```

```

def _jaccard(a_rows, a_cols, b_rows, b_cols):
    """Jaccard coefficient on the elements of the two biclusters."""
    intersection = ((a_rows * b_rows).sum() *
                    (a_cols * b_cols).sum())

    a_size = a_rows.sum() * a_cols.sum()
    b_size = b_rows.sum() * b_cols.sum()

    return intersection / (a_size + b_size - intersection)

def _pairwise_similarity(a, b, similarity):
    """Computes pairwise similarity matrix.

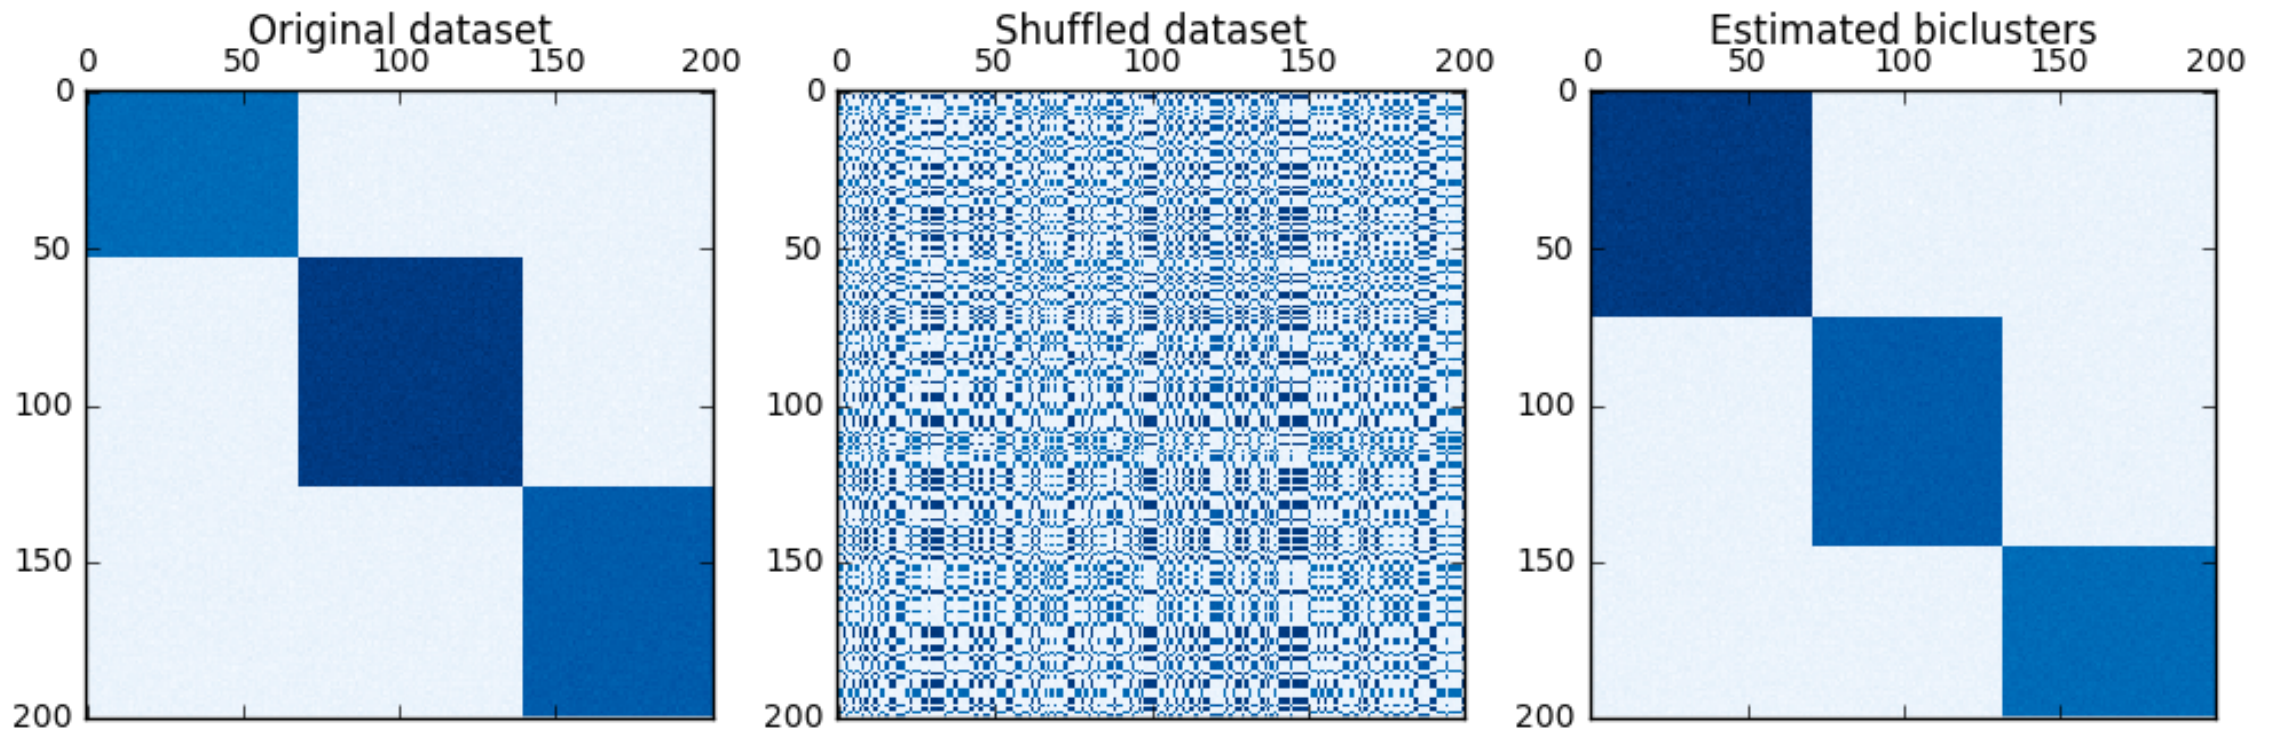
    result[i, j] is the Jaccard coefficient of a's bicluster i and b's
    bicluster j.

    """
    a_rows, a_cols, b_rows, b_cols = _check_rows_and_columns(a, b)    # unpacks rows and columns
    n_a = a_rows.shape[0]
    n_b = b_rows.shape[0]
    result = np.array(list(list(similarity(a_rows[i], a_cols[i],
                                           b_rows[j], b_cols[j])
                                   for j in range(n_b))
                           for i in range(n_a)))

    return result

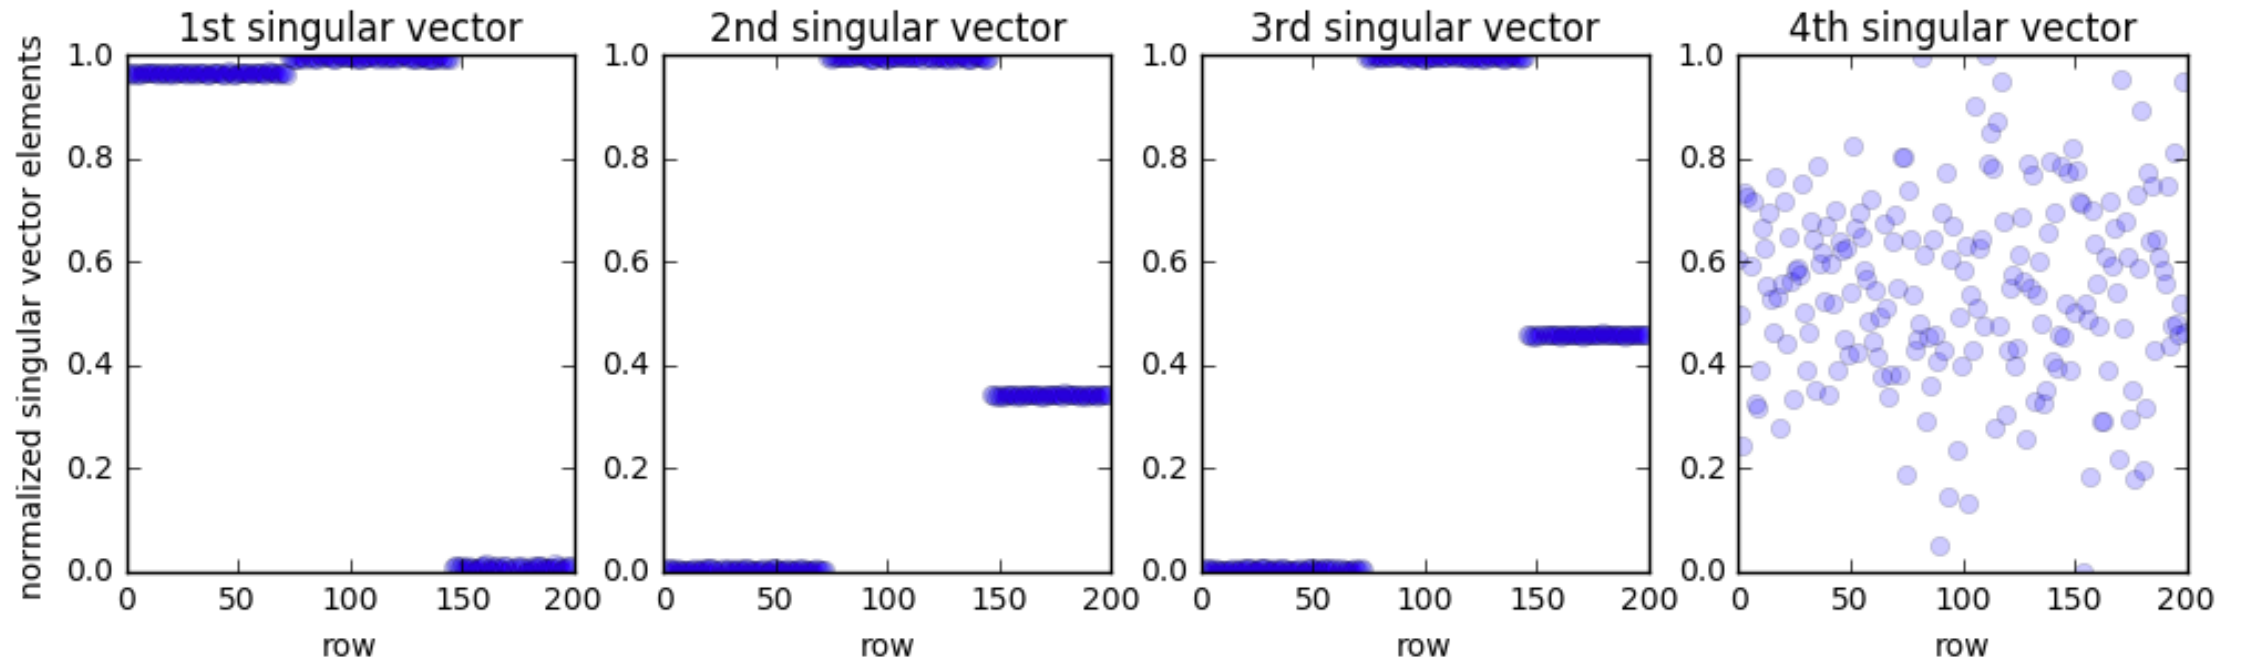
```

Example




- Recovery of hidden co-clusters by spectral co-clustering

Example: singular vectors



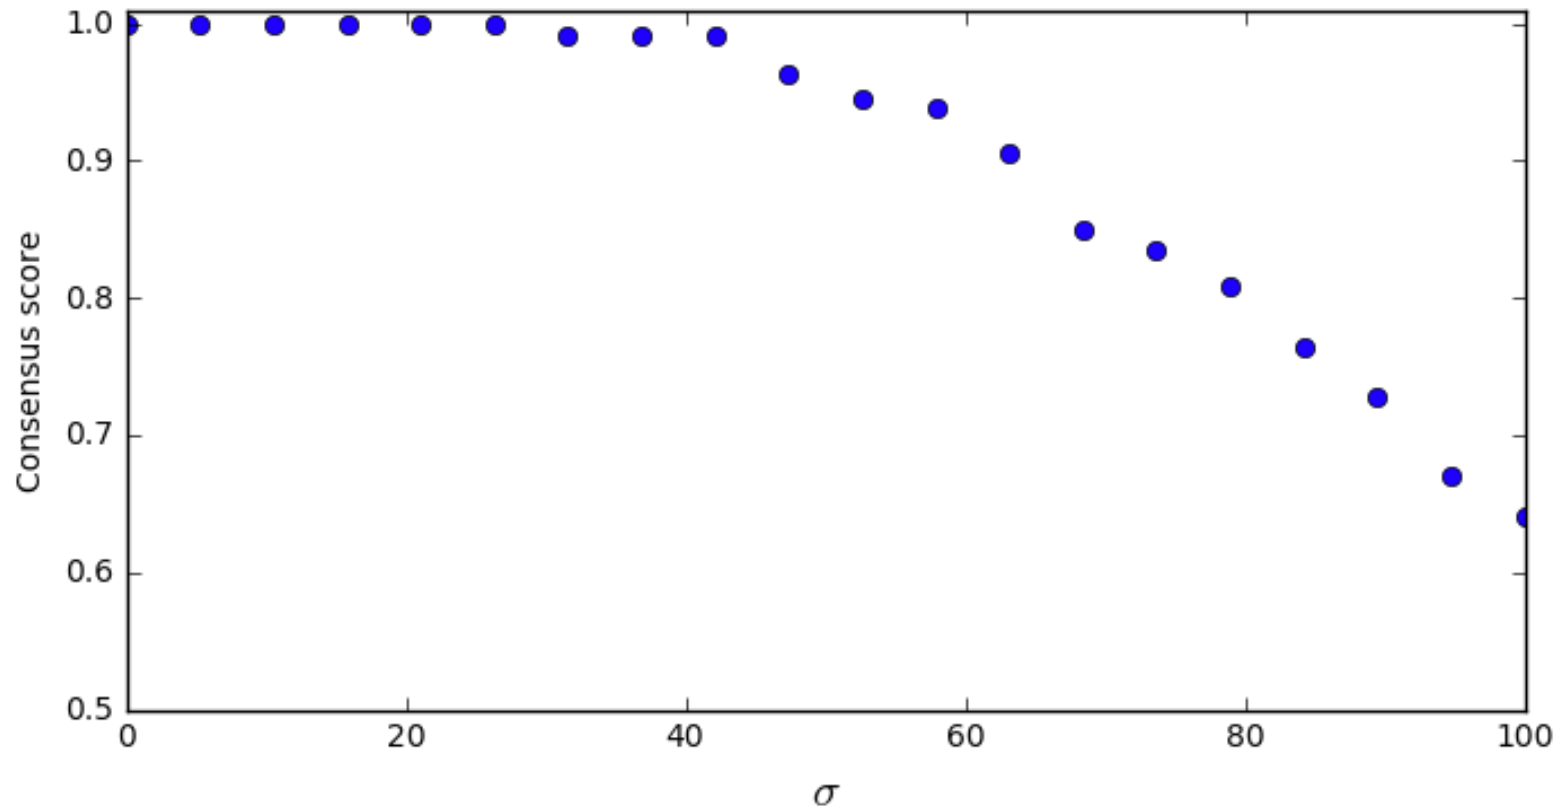
- Clusters are clearly indicated by the 2nd and 3rd singular vectors

Robustness to noise



- Input matrix corrupted by noise with varying variance σ

Robustness to noise (cont'd)



- Consensus score vs standard deviation of noise

References

- J. E. Atkins, E. G. Boman and B. Hendrickson, A Spectral Algorithm for Seriation and the Consecutive Ones Problem, SIAM J. Computing, Vol 28, No 1, pp. 297-310, 1998
- Y. Benjamini, Opening the Box of a Boxplot, The American Statistician, Vol 42, No 4, pp 257-262, 1988
- I. S. Dhillon, Co-clustering documents and words using bipartite spectral graph partitioning, Proc. of ACM KDD, pp. 269-274, 2001
- M. Eisen, P. Spellman, P. Brown and D. Botstein, Cluster analysis and display and display of genome-wide expression patterns, PNAS, Vol 95, pp 14863-68, 1998
- M. Friendly, Corrgrams: exploratory displays for correlation matrices, The American Statistician, Vol 56, pp. 316-324, 2002
- Y. Kluger, R. Basri, J. T. Chang, M. Gerstein, Spectral Biclustering of Microarray Data: Coclustering Genes and Conditions. Genome Research, 13(4), pp. 703-716 , 2003
- L. Wilkinson and M. Friendly, The history of the cluster heat map, The American Statistician, Vol 63, No 2, pp 179-184, 2009